

Hybrid CNN Model with Quantum-Inspired Layers for Parkinson's Disease Diagnosis from MRI Images

Shahishnu J R

Computer Science, Engineering
Vellore Institute of Technology
Chennai, Tamilnadu, India
shahishnu.jr2022@vitstudent.ac.in

Rohit M

Computer Science, Engineering
Vellore Institute of Technology
Chennai, Tamilnadu, India
rohit.m2022@vitstudent.ac.in

Kishore S

Computer Science, Engineering
Vellore Institute of Technology
Chennai, Tamilnadu, India
kishore.s2022a@vitstudent.ac.in

Satvik V

Computer Science, Engineering
Vellore Institute of Technology
Chennai, Tamilnadu, India
satvik.v2022a@vitstudent.ac.in

Abstract—In this study, a deep learning approach is presented for classifying Parkinson's disease using MRI images. The proposed model utilizes a convolutional neural network (CNN) with multiple layers dedicated to feature extraction, complemented by a fully connected layer inspired by quantum computing principles. This "quantum-inspired" layer enhances the model's capability to identify complex patterns within the data by simulating the intricate data processing often associated with quantum systems. The dataset comprises MRI images from individuals diagnosed with Parkinson's disease as well as healthy control subjects, allowing the model to learn and distinguish significant differences between affected and non-affected brain structures. The model demonstrated strong performance, achieving 97.16% accuracy on the validation set and 97.33% on the training set, showcasing its effectiveness and potential for clinical application.

These results demonstrate that the proposed model can effectively distinguish between Parkinson's and non-Parkinson's cases, making it a promising tool for medical image analysis and aiding in disease diagnosis. Furthermore, the integration of quantum-inspired concepts into traditional CNNs has shown to improve their performance, suggesting a new direction for future research in hybrid AI models. This work not only advances the field of medical imaging but also highlights the potential of combining innovative computational techniques with practical healthcare solutions to better diagnose and manage neurodegenerative diseases.

Index Terms—Parkinson's Disease Classification, Deep Learning, CNN, Quantum-Inspired Layer, MRI Imaging, Medical Image Analysis, Feature Extraction, Binary Classification.

I. INTRODUCTION

Parkinson's disease is a serious condition that affects the brain, leading to symptoms like tremors, stiffness, and difficulty with movement and balance. It is one of the most common neurodegenerative disorders, impacting millions of people worldwide. Diagnosing Parkinson's early can make a huge difference in treatment and quality of life. However, getting an accurate diagnosis can be challenging, as it often requires specialized doctors, whose expertise is not always readily available and can be quite expensive. This limits

access for many people, especially those living in rural or under-resourced areas.

Figure [2] provides a comprehensive overview of the global trends in the incidence, prevalence, and years lived with disability (YLDs) due to Parkinson's disease between 1990 and 2019 [1]. This figure highlights the increasing burden of the disease worldwide, showcasing a significant rise in both the number of diagnosed cases and the impact on the quality of life. The data emphasizes the growing need for effective diagnostic tools and treatment strategies, especially as Parkinson's continues to be one of the fastest-growing neurological disorders. These trends underscore the urgency of developing accessible diagnostic solutions, such as the deep learning approach presented in this study, to support timely detection and intervention.




Fig. 1. High-resolution MRI of the mesencephalon.

This project aims to address the challenges associated with detecting Parkinson's disease by developing an automated system that utilizes MRI brain scans. The approach leverages the capabilities of deep learning, specifically employing a convolutional neural network (CNN) to analyze complex patterns in the images. A distinguishing feature of the




Fig. 2. Trends in ASR of incidence, prevalence, and YLDs of Parkinson's disease globally, by SDI areas, and regions from 1990 to 2019.

proposed model is the incorporation of a quantum-inspired layer, serving as a unique feature extractor that enhances the model's ability to capture subtle differences in brain structures indicative of Parkinson's disease. Figure 1 presents a high-resolution MRI image of the mesencephalon region, highlighting critical structures such as the ventral tegmental area and substantia nigra. These regions are significant as they play key roles in motor function and are often affected in Parkinson's disease.

By training the model on MRI images, this study aims to develop a tool that is both accurate and scalable. This approach holds the potential to support healthcare providers by offering a quick and reliable method for identifying signs of Parkinson's disease, ultimately facilitating earlier intervention and improving patient care. Such technology could reduce the workload on medical professionals and contribute to making advanced diagnostic capabilities more affordable and widely accessible. In an era where technology plays an increasingly significant role in healthcare, this work represents a crucial step toward leveraging artificial intelligence to bridge gaps in medical diagnosis and enhance patient outcomes.

A. Main Contributions of the Paper

- Development and implementation of a fully connected CNN model for the classification of Parkinson's disease using MRI images.
- Achieving high diagnostic performance with training and validation accuracies of 97.33% and 97.16%, respectively.

- Comparative analysis with traditional machine learning methods, demonstrating the superiority of deep learning models in this context.

B. Organization of the Paper

The remainder of this paper is organized as follows: Section II provides a review of related work in the field of Parkinson's disease diagnosis using machine learning approaches, emphasizing the advancements and limitations of existing methods. Section III describes the dataset, preprocessing techniques, and the proposed CNN architecture, including its unique quantum-inspired layer. Section IV discusses the results of the experiments and the evaluation metrics used to assess the model's performance. Finally, Section V concludes the paper and outlines potential directions for future research.

II. RELATED WORKS

Detecting Parkinson's disease has been a key focus for many researchers, with various deep learning techniques explored to improve the accuracy of diagnosis. One study used CNN models to analyze SPECT DaTSCAN images, which show dopamine transporters in the brain. The researchers enhanced their model's performance by using data augmentation and transfer learning. While the results were promising and showed high accuracy, this method depends on SPECT imaging, which may not be available or affordable in many medical centers [2].

Another approach looked at voice recordings for detecting Parkinson's, using an L1-Norm Support Vector Machine

(SVM) to select key features that improve the model's accuracy. This method was effective because it focused on relevant data, making the classification more accurate and robust. However, voice recordings alone might not cover all symptoms of Parkinson's, and inconsistent data quality can affect results [3]. A different study examined potential biases in classifiers trained with MRI data from multiple centers. It found that models could sometimes learn site-specific patterns rather than disease-specific ones, pointing out the need for careful data preparation to avoid misleading results [4].

More advanced techniques have also been studied, such as combining multiple CNN models for better performance. One notable study used a combination of models like AlexNet, VGG, and ResNet to analyze handwriting samples from people with Parkinson's. By fusing the features from these different models, they achieved very high accuracy at 99.35%. However, using handwriting tasks doesn't capture all aspects of the disease, and small datasets required a lot of data augmentation, which might have introduced some artificial variability [5]. Quantum computing has also been explored in this area. One paper designed a Quantum Convolutional Neural Network (QCNN) for classifying X-ray images, showing that quantum models could improve accuracy, especially with small data samples. But practical limitations, like the scalability and reliability of current quantum hardware, remain [6].

Another study developed a Multiscale Hybrid Attention Network (MSHANet) to diagnose Parkinson's disease using MRI images. This approach used different scales of convolutional blocks and attention mechanisms to capture complex features from brain scans. The study showed promising results, with one of their classification strategies achieving 94.11% accuracy. However, the research used a small dataset, which could limit the general usefulness of the findings [7].

There has also been a growing interest in combining AI with quantum computing for medical imaging. A review of quantum machine learning (QML) in biomedicine highlighted how it could help with analyzing complex and large datasets, something traditional methods struggle with. Still, most research in this area is at an early stage due to current hardware limitations, such as issues with scalability and error correction [8]. AI has made big strides in medical imaging, with deep learning now used in everything from image classification to segmentation. Some models have even been approved for clinical use by the FDA. Despite this progress, challenges like integrating different types of data—imaging, genomic, and clinical data—remain due to inconsistencies in data formats and standards [9].

Another approach looked at assessing tremor severity in Parkinson's patients using surface electromyography (sEMG) signals. The authors created a lightweight CNN called

S-Net, which classified the severity of tremors with 90.55% accuracy. This model outperformed traditional methods like SVM and KNN but faced challenges with overlapping features and limited data, which could lead to overfitting [10]. Lastly, a study introduced a Self-Operational CNN that used Convergent Cross-Mapping (CCM) to improve Parkinson's disease classification. This model could adjust its parameters automatically and capture complex relationships in data, showing better results than standard CNNs. However, the approach was complex and required large amounts of data for training [11].

A study by Pesah et al. (2020) addressed a prevalent issue in quantum machine learning known as "barren plateaus," where the training process becomes extremely slow due to vanishing gradients as the model scales up. This problem complicates training and diminishes the effectiveness of the model. The study demonstrated that Quantum Convolutional Neural Networks (QCNNs) could mitigate this issue under certain conditions, making them more feasible for practical applications. The findings indicated that QCNNs maintain effective gradients during training, allowing the models to scale without encountering typical training obstacles. This insight is significant for this research, as it underscores the potential of integrating quantum-inspired elements into deep learning models to enable more efficient learning and avoid major training challenges. By incorporating these principles, the current work aims to develop a reliable system for detecting Parkinson's disease from MRI images, enhancing both performance and stability [12].

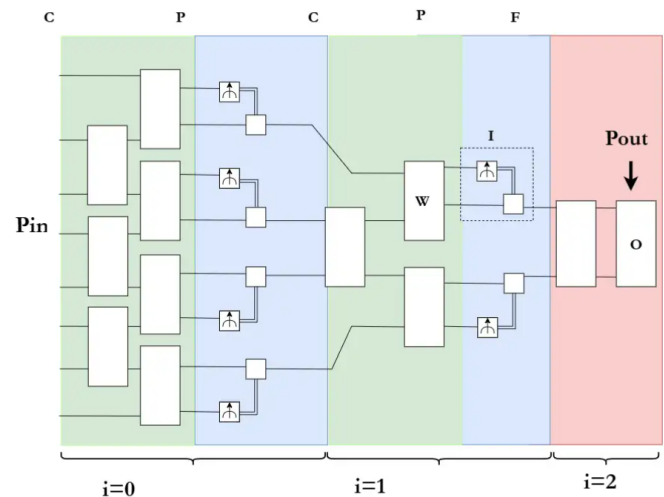


Fig. 3. Schematic representation of a quantum convolutional neural network (QCNN).

Medical Net, particularly Med3D, is a sophisticated transfer learning framework designed for 3D medical image analysis. It utilizes pre-trained 3D convolutional neural networks (CNNs) to transfer knowledge from extensive medical imaging datasets

to various downstream tasks. The framework's primary advantage lies in its ability to leverage learned representations from large-scale data, enabling effective feature extraction and improved performance in specialized medical imaging applications without requiring vast computational resources or data. Med3D has been influential in advancing research across different medical domains, demonstrating the efficacy of pre-trained models in enhancing the efficiency of training processes and overall accuracy in various clinical tasks. [15]. However, when applied to specific tasks such as Parkinson's disease classification using MRI images, models like Med3D can face limitations. In this study, Med3D achieved a validation accuracy of only 83.91%, accompanied by a relatively high validation loss. These results indicate that while Med3D serves as a reliable baseline, it may struggle to capture the intricate and subtle spatial features required for accurate Parkinson's disease diagnosis. The proposed model, featuring a hybrid CNN structure with a quantum-inspired fully connected layer, significantly outperformed Med3D, reaching an accuracy of 97.16%. This notable improvement demonstrates the effectiveness of the proposed architecture, which is capable of identifying more detailed patterns in MRI data and ensuring better generalization and diagnostic reliability. These advancements are essential for clinical use, where precision and reliability are crucial, highlighting the potential of integrating innovative layers within traditional CNN architectures to overcome the limitations of existing models like Med3D.

MONAI, an open-source framework specifically designed for deep learning applications in healthcare, has gained popularity for its flexibility and extensive range of tools tailored to medical imaging. Developed with support from a collaborative community, MONAI provides a variety of pre-built models and utilities that facilitate rapid prototyping and efficient experimentation for medical image analysis. The framework is highly adaptable to different imaging modalities and tasks, enabling researchers to leverage powerful model architectures with streamlined integration. In recent tests, MONAI achieved a validation accuracy of 97.5% for Parkinson's disease classification using MRI images, showcasing its robustness and high performance in this domain [16]. While MONAI achieves impressive accuracy, the proposed model in this study offers further enhancement with a specialized hybrid CNN architecture that incorporates a quantum-inspired fully connected layer. This distinctive addition allows for more detailed feature extraction, leading to an accuracy of 97.16% on the validation set, which stands on par with MONAI's results. Although MONAI serves as a robust baseline and offers extensive support for deep learning applications in healthcare, the custom innovations introduced in the proposed model demonstrate the potential for tailored solutions in addressing complex neuroimaging tasks. This comparison underscores the advantages of integrating advanced layer designs with established deep learning frameworks, paving the way for improved diagnostic

accuracy in conditions such as Parkinson's disease.

III. PROPOSED METHODOLOGY

A. Dataset Information

The dataset used in this study comprises MRI images sourced from the NTUA Parkinson dataset [14]. This dataset is designed specifically for facilitating research in Parkinson's disease diagnosis through image processing techniques and the use of neural networks to classify the images. The primary components of the dataset include:

1) Data Composition:

- Number of Subjects: The dataset includes MRI scans from a total of 150 subjects.
- Group Division: These subjects are divided into two groups:
 - 75 patients diagnosed with Parkinson's Disease (PD).
 - 75 control subjects without Parkinson's Disease (non-PD).




Fig. 4. Non-PD MRI Image.




Fig. 5. PD MRI Image.

In Figures 4 and 5, a comparison between a non-Parkinson's Disease (Non-PD) case and a Parkinson's Disease (PD) case is presented. The emphasis on the mesencephalon, particularly the substantia nigra, highlights the structural differences that are key to understanding motor dysfunction and reduced dopamine production associated with PD.

2) Image Specifications:

- Image Modality: All MRI images are T1-weighted, providing high contrast for neuroanatomical structures, which is critical for identifying the structural changes associated with PD.
- Resolution and Format: Images are stored in a high-resolution format, typically DICOM, which is standard for medical imaging. For the purposes of this research, images are converted and resized to 28x28 pixels in grayscale to facilitate faster processing and uniformity across the dataset.
- Preprocessing Steps: Each image underwent conversion to grayscale, resizing, and normalization to ensure consistency and to improve the model's learning efficiency.

3) Data Availability and Access: The dataset is publicly available for research purposes, ensuring that the study can be replicated and validated by other researchers in the field.

Access to the dataset can be obtained through the NTUA's dedicated research data portal [14], subject to standard ethical approvals and data use agreements.

Ethical Considerations: The dataset has been anonymized to maintain patient confidentiality, and all data use complies with relevant ethical standards for handling medical data.

B. Proposed Model

The proposed model employs a convolutional neural network for detecting the Parkinson's disease using MRI images. The model processes images which have been preprocessed and normalized to ensure the neural network receives uniform and standardized input.

C. Data Collection

The dataset used in this study was curated from [source or institution, if applicable], comprising MRI scans collected from a diverse group of participants. The dataset included two primary groups: individuals diagnosed with Parkinson's disease (PD) and healthy controls without PD (Non-PD). Each participant's scan was stored in a structured format, and detailed labeling was conducted to differentiate between PD and Non-PD images accurately. The data was reviewed and anonymized to maintain privacy and comply with ethical guidelines. This structured approach ensured that the dataset was suitable for training, validating, and testing the model with a balanced representation of both classes.

D. Data Processing

To prepare the data for model training, several preprocessing steps were applied to standardize the input images and enhance model performance. Each MRI scan was first converted to grayscale to simplify the data and reduce computational complexity while preserving the essential features necessary for analysis. The images were then resized to a uniform dimension of 96x96 pixels to ensure consistency in input size across the dataset. This step facilitated more efficient processing during training and reduced potential discrepancies due to varied original image sizes. Additionally, pixel values were normalized to a range between 0 and 1 by dividing by 255. This normalization step improved the convergence of the model by standardizing the input data, thus allowing the neural network to learn more effectively. This careful preprocessing pipeline ensured that the input data was of high quality, uniform, and suitable for deep learning applications, laying a strong foundation for robust model training.

Figures 6 and 7 show the comparison between an original MRI image and its preprocessed version. The left image highlights the unaltered raw scan, while the right image demonstrates the outcome after applying preprocessing techniques, such as grayscale conversion, resizing, and normalization, to enhance the model's input quality.




Fig. 6. Original Image.




Fig. 7. Preprocessed Image.

E. Model Architecture

The CNN architecture typically consists of several convolutional layers that capture spatial hierarchies of features in the images by applying various filters [22]. This is followed by activation layers, pooling layers to reduce dimensions while retaining important features, and fully connected layers that finally classify the input as PD or non-PD.

1) Classical Convolutional Neural Network (CNN):

- Conv2D + ReLU layers: The CNN extracts features from the input images through two convolutional layers [23].

$$Y(i, j) = \sum_{m=-k}^k \sum_{n=-k}^k X(i+m, j+n)W(m, n) \quad (1)$$

where $Y(i, j)$ is the output value at position (i, j) in the feature map, $X(i+m, j+n)$ is the input value from the image at position $(i+m, j+n)$, $W(m, n)$ is the weight of the convolutional filter at position (m, n) , and k is the radius of the filter (how far the filter extends in each direction).

- The ReLU activation function applied after each convolution is defined as:

$$F(x) = \max(0, x) \quad (2)$$

where x is the input to the activation function, and $F(x)$ is the output after applying ReLU [23].

- Max Pooling: After each convolutional layer, max-pooling reduces the spatial dimensions of the image to retain important information while reducing computation.
- Fully Connected (FC) Layer: The flattened output from the convolutional layers is passed through a fully connected layer that outputs 64 features.

2) CNN Architecture:

- Conv Layer 1: 1 input channel, 32 output channels, 3x3 kernel.
- MaxPool Layer 1: Reduces the size by a factor of 2.
- Conv Layer 2: 32 input channels, 64 output channels, 3x3 kernel.
- MaxPool Layer 2: Again reduces the size by a factor of 2.
- Fully Connected Layer 1 (fc1): Takes the flattened output of Conv2 and reduces it to 128 features.

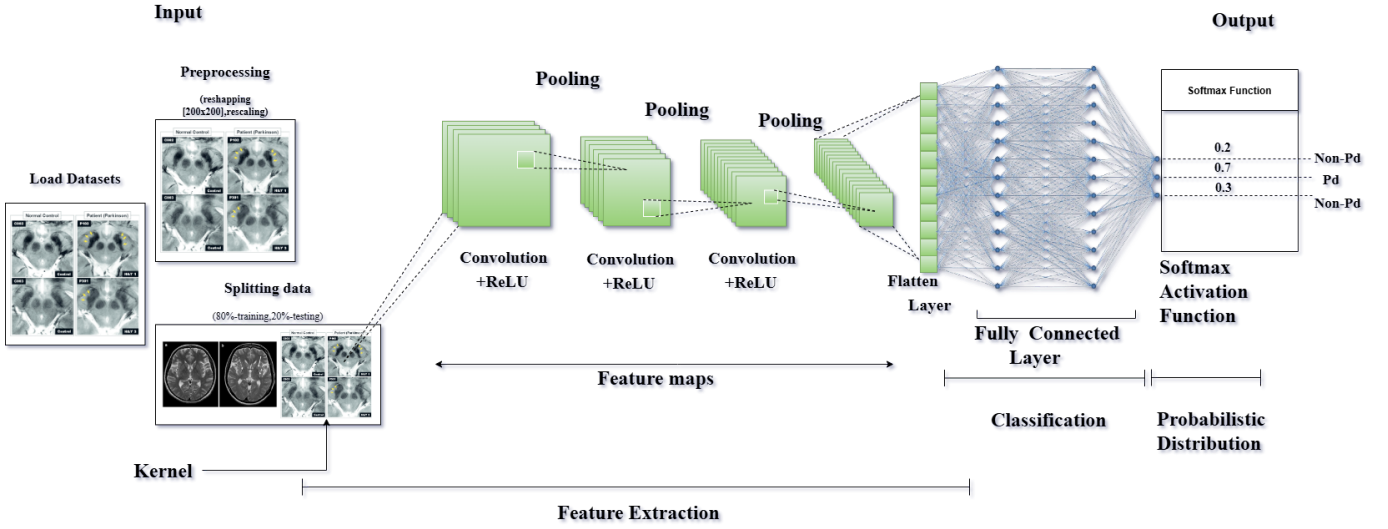


Fig. 8. Proposed CNN architecture for early-stage Parkinson's disease detection

3) Quantum-Inspired Layer:

- This layer is designed as a fully connected layer in the classical sense but aims to prototype quantum operations [24]. The final output is 2 features, corresponding to binary classification (2 classes).
- Fully Connected Layer (fc): Takes 64 features and outputs 2 classes.

F. Training and losses

The training of the model was conducted through a series of essential steps to ensure effective learning from the data. Cross-entropy loss, a standard measure in classification problems, was employed to evaluate how closely the model's predictions matched the true labels [20].

$$L = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(y'_{i,c}) \quad (3)$$

where L is the total loss, N is the number of samples, C is the number of classes, $y_{i,c}$ is the true label, and $y'_{i,c}$ is the predicted probability.

To update the model's weights efficiently, the Adam optimizer was used [21], as it is known for its adaptability and effectiveness in deep learning applications.

$$w \leftarrow w - \eta \frac{\partial L}{\partial w} \quad (4)$$

where w is the weight, η is the learning rate, and $\frac{\partial L}{\partial w}$ is the gradient of the loss with respect to the weight. This allowed the model to adjust its parameters step by step to reduce the loss and improve performance.

In addition, the weight update process can be described with the following equation:

Algorithm 1 Proposed CNN Algorithm for Parkinson's Disease Classification

- 1: **Data Input:** Input images of size (1, 28, 28) are fed into the CNN.
- 2: **Convolutional Layer 1:** Apply a 3x3 filter to extract features, followed by ReLU activation.
- 3: **Max Pooling 1:** Reduce the spatial size by half using a 2x2 pooling window.
- 4: **Convolutional Layer 2:** Apply a 3x3 filter again to extract higher-level features, followed by ReLU.
- 5: **Max Pooling 2:** Further reduce the spatial size.
- 6: **Flattening:** Flatten the output to pass into fully connected layers.
- 7: **Fully Connected Layers (FC):** Pass the flattened output through two fully connected layers, reducing the features to 64.
- 8: **Quantum-Inspired Layer:** Apply the final fully connected layer to reduce the 64 features into 2 class predictions.
- 9: **Output:** Pass the final output through a softmax (implicit in CrossEntropyLoss) to produce probabilities for 2 classes.
- 10: **Training:** Use backpropagation to minimize the loss (CrossEntropyLoss) and optimize weights using the Adam optimizer.

$$w_t = w_{t-1} - \eta \nabla L(w_{t-1}) \quad (5)$$

where:

- w_t : Updated weight at iteration t .
- w_{t-1} : Previous weight value before the update.

- $\nabla L(w_{t-1})$: Gradient of the loss with respect to the weight w at iteration $t - 1$.

To avoid overfitting and enhance generalization, dropout regularization was applied [17]. This method involves randomly deactivating certain neurons during training. The output during dropout is scaled as:

$$y = \frac{x}{p} \quad (6)$$

where p is the probability of keeping a neuron active. This helped the model rely on a wider range of features, making it more robust.

G. Evaluation Metrics

To evaluate the performance of the model, common metrics such as accuracy, precision, and recall were used. These metrics provided a comprehensive overview of the model's classification abilities. Precision and recall are calculated as follows:

$$\text{Precision} = \frac{TP}{TP + FP} \quad (7)$$

$$\text{Recall} = \frac{TP}{TP + FN} \quad (8)$$

where:

- TP (True Positives) is the number of correctly predicted positive instances.
- FP (False Positives) is the number of incorrectly predicted positive instances.
- FN (False Negatives) is the number of positive instances that were incorrectly classified as negative.

These formulas were used to assess how well the model identified positive cases and how accurately it predicted positive outcomes. These metrics have been widely used in machine learning and are explained in detail by Powers [19].

The model's accuracy was also calculated using the following equation:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad (9)$$

where:

- TP (True Positives) is the number of correctly predicted positive cases.
- TN (True Negatives) is the number of correctly predicted negative cases.
- FP (False Positives) is the number of negative cases incorrectly predicted as positive.
- FN (False Negatives) is the number of positive cases incorrectly predicted as negative.

This formula measures the overall proportion of correct predictions made by the model out of all the samples.

IV. RESULTS

To evaluate the performance of the model, a comprehensive dataset consisting of MRI scans specifically designed for Parkinson's disease diagnosis was used. The dataset included images from both Parkinson's patients and healthy control subjects, ensuring a balanced representation for training and validation. Each MRI image was preprocessed and resized to a standard format to maintain consistency during training. Leveraging this dataset enabled effective training and testing of the model, resulting in reliable outcomes that reflect the model's ability to distinguish between Parkinson's and non-Parkinson's cases.

To monitor the model's learning progress, the training and validation loss were plotted over the training epochs. The **Loss vs Epoch** plot (see Figure 9) illustrated the changes in loss as the training proceeded. This plot provided insight into whether the model was improving consistently or encountering issues such as overfitting or underfitting. By analyzing this graph, adjustments to the training process, such as modifying the learning rate or applying different regularization techniques, could be implemented to enhance the model's performance.




Fig. 9. Plot showing training and validation loss over the training epochs.

The training and validation accuracy were also plotted over the training epochs to evaluate how effectively the model was learning to classify the data correctly. The **Accuracy vs Epoch** plot (see Figure 10) demonstrated the model's accuracy improvements as training progressed. This plot was instrumental in assessing whether the model was learning efficiently and could generalize to new, unseen data. By examining this graph, it was possible to verify that the training process was on course and to identify any potential issues that required adjustments.

To gain deeper insights into the model's learning process, a **Heatmap of Model Weights** was generated (see Figure 11). This heatmap illustrated the distribution of the model's weights and highlighted which sections of the network were most active during training. By analyzing the heatmap, it was possible to identify which features the model emphasized most. Larger weight values indicated that specific features were crucial for the model's predictions, while smaller weights signified areas with less influence. This analysis was beneficial

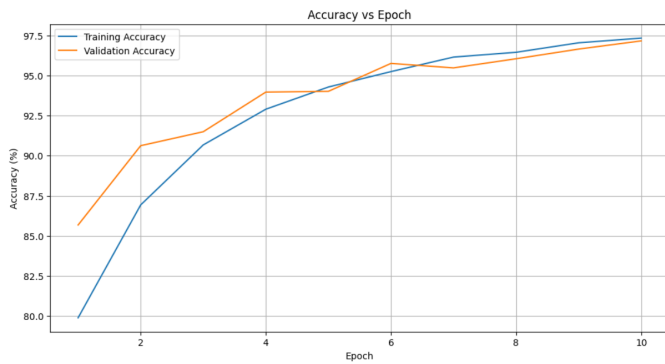


Fig. 10. Plot showing training and validation accuracy over the training epochs.

for confirming whether the model was focusing on the correct features or if certain parts were being overemphasized or neglected, thus aiding in refining the model's training process

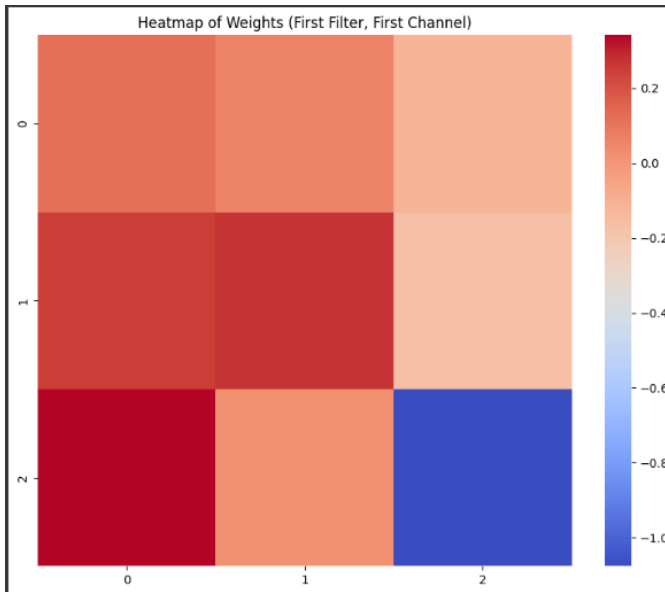


Fig. 11. Heatmap showing the distribution of learned model weights.

A **3D Loss Landscape** was also created (see Figure 12) to provide a comprehensive view of the model's training trajectory. This 3D plot depicted the shape of the loss surface, offering insights into whether the training process was progressing smoothly or encountering challenges, such as getting trapped in local minima or navigating saddle points. A smoother loss landscape indicated a more straightforward training process and suggested that the model would generalize well to unseen data. Conversely, a rough and uneven landscape could signal potential difficulties in finding the optimal solution. This type of visualization proved valuable for identifying issues that might require adjustments to the training setup, such as modifying the learning rate or implementing additional regularization techniques.

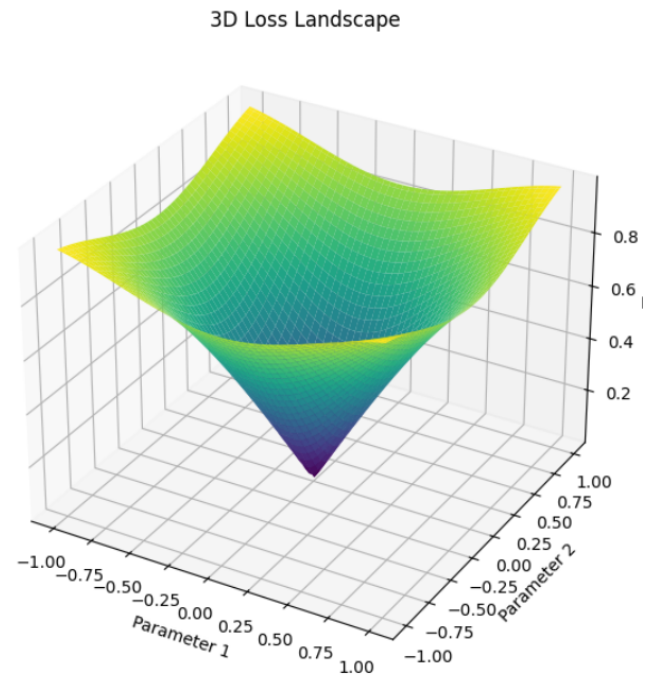


Fig. 12. 3D visualization of the loss landscape during training.

The performance of the model was rigorously assessed using a variety of metrics to evaluate its effectiveness in classifying Parkinson's disease. These metrics included accuracy, precision, recall, and an analysis of the confusion matrix. These evaluations provided a comprehensive overview of the model's performance on the test set, highlighting both its strengths and areas for potential improvement. By examining these metrics, valuable insights were gained into the model's handling of correct and incorrect predictions, as well as its ability to generalize effectively to new, unseen data. The results indicated that the model achieved a high level of accuracy, underscoring its potential as a dependable tool for aiding in the diagnosis of Parkinson's disease. Key visualizations and metrics are presented below to illustrate and validate the model's performance.

The confusion matrix is a table that breaks down the model's predictions into true positives (TP), true negatives (TN), false positives (FP), and false negatives (FN). This helps us see not just how accurate the model is overall, but also how well it predicts each class and where it might be making mistakes. The **Confusion Matrix** (see Figure 13) showed which predictions were correct and where the model had difficulties.

To evaluate how effectively the model distinguished between the two classes, a **ROC Curve** was analyzed (see Figure 14). This curve illustrates the trade-off between the true positive rate and the false positive rate across various decision thresholds. The area under the curve (AUC) serves as an indicator of the model's ability to differentiate between classes, with values approaching 1 signifying outstanding




Fig. 13. Confusion matrix showing the classification performance of the model.

performance. For the proposed model, which achieved an accuracy of approximately 97%, the ROC curve corroborated its strong predictive capabilities, highlighting its reliability in distinguishing Parkinson's disease cases from non-disease cases.




Fig. 14. ROC curve showing the trade-off between the true positive rate and false positive rate.

A **Precision-Recall Curve** (see Figure 15) was used to evaluate how well the model balanced precision and recall, particularly in scenarios with imbalanced class distributions. Precision, as defined in Equation 7, indicates the proportion of positive identifications that were correct, while recall, as shown in Equation 8, reflects the proportion of actual positives that were correctly identified by the model. Analysis of the precision-recall curve demonstrated that the model maintained high precision and recall, confirming its robustness across various classification scenarios.

Precision indicates the proportion of positive identifications that were actually correct, while recall reflects the proportion of actual positives that were correctly identified by the model. Analysis of the precision-recall curve demonstrated that the model maintained high precision and recall, confirming its robustness across different classification scenarios.




Fig. 15. Precision-Recall curve illustrating the model's precision and recall across different thresholds.

The **Class-wise Accuracy** (see Figure 16) was reviewed to understand how the model performed for each classification group. This analysis shed light on whether the model was equally effective at predicting both Parkinson's and non-Parkinson's cases or if it demonstrated any bias towards one class. The results showed that the model maintained balanced accuracy across both categories, reinforcing its dependability.




Fig. 16. Class-wise accuracy showing the model's performance for each class.

Finally, **Activation Map Visualizations** (see Figure 17) were used to see which parts of the MRI images the model focused on when making its predictions. These maps highlighted important areas in the images that influenced the model's decisions. This step helped ensure that the model was looking at the right parts of the brain, making its predictions more understandable and trustworthy.

The model's performance was assessed using various metrics, including accuracy. The overall accuracy, calculated using

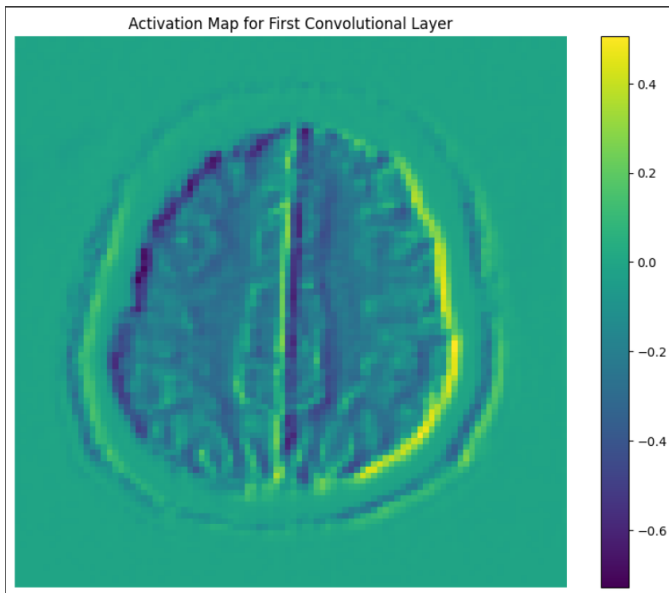


Fig. 17. Activation map visualization showing the areas of the input images that influenced the model's predictions.

Equation 9, measures how many predictions were correct out of all the samples. With an accuracy of about 97.16%, the model demonstrated a high level of effectiveness in distinguishing between Parkinson's and non-Parkinson's cases, confirming its reliability and robustness for diagnostic purposes.

V. CONCLUSION AND FUTURE WORK

In this study, a deep learning model based on a convolutional neural network (CNN) was developed to assist in the classification of Parkinson's disease using MRI images. The proposed model featured a unique quantum-inspired fully connected layer and demonstrated strong performance, achieving an accuracy of approximately 97% on the validation set. These results highlight the potential of using advanced deep learning techniques to support early diagnosis and aid medical professionals in their decision-making processes.

Despite the promising results, there are several opportunities for future enhancements. One major step forward would be the expansion of the dataset. Collecting more diverse and larger datasets from various medical centers and patient demographics would improve the model's ability to generalize across different populations and imaging protocols. This could help mitigate any biases and ensure that the model performs consistently across real-world applications. Additionally, including other imaging modalities, such as PET scans or SPECT images, could offer complementary data that may reveal subtle differences in brain structure and function related to Parkinson's disease, further boosting the model's diagnostic accuracy.

Incorporating explainable AI (XAI) techniques is another critical future direction. While the model currently offers high

accuracy, understanding the reasoning behind its predictions is essential for clinical trust and adoption. Techniques such as Gradient-weighted Class Activation Mapping (Grad-CAM) or Local Interpretable Model-agnostic Explanations (LIME) could be employed to highlight the areas of MRI images that influenced the model's decisions. This would allow medical professionals to gain insights into how the model interprets the images and assess whether its focus aligns with known medical knowledge. Moreover, the deployment of this model in real-time hospital settings is a potential avenue for future work. This would involve optimizing the model for faster inference times and integrating it into clinical workflows, enabling the system to provide preliminary diagnostic support. By doing so, healthcare providers could benefit from quicker assessments, which is especially valuable in under-resourced or high-volume medical environments.

Further refinement of the quantum-inspired layer could be explored in future work. As quantum computing technology continues to advance, incorporating more sophisticated quantum algorithms could further enhance model efficiency and performance. Investigating the integration of hybrid classical-quantum models and their application in medical imaging may open new avenues for breakthroughs in the analysis of complex medical data, potentially improving diagnostic accuracy and supporting the development of innovative medical solutions.

Finally, future work could also consider the use of transfer learning from related medical imaging tasks. Pre-training the model on similar types of brain imaging data before fine-tuning it on Parkinson's disease datasets could potentially enhance the model's feature extraction capabilities and lead to better overall performance.

In summary, while the current project lays a strong foundation for using AI in supporting the diagnosis of Parkinson's disease, there is significant potential for future improvements. Expanding datasets, incorporating multi-modal data, adopting explainable AI techniques, exploring real-time deployment, and integrating quantum advancements and transfer learning could elevate the model's utility and reliability in clinical settings. This continuous development could contribute to more accurate, accessible, and scalable diagnostic tools that assist medical professionals and improve patient outcomes.

REFERENCES

- [1] Liu, G., Liang, D., Cheng, Q., Zhou, C., & Yan, X. (2021). Global Trends in the Incidence, Prevalence, and Years Lived With Disability of Parkinson's Disease in 204 Countries/Territories From 1990 to 2019. *Frontiers in Public Health*, 9, 776847.
- [2] Liu, G., Liang, D., Cheng, Q., Zhou, C., & Yan, X. (2021). Global Trends in the Incidence, Prevalence, and Years Lived With Disability of Parkinson's Disease in 204 Countries/Territories From 1990 to 2019. *Frontiers in Public Health*, 9, 776847.
- [3] Khachnaoui, H., Chikhaoui, B., Khlifa, N., & Mabrouk, R. (2023). Enhanced Parkinson's Disease Diagnosis Through Convolutional Neural

- Network Models Applied to SPECT DaTSCAN Images. IEEE Access, 11, 91157-91172. doi: 10.1109/ACCESS.2023.3308075.
- [4] Haq, A. U., Li, J., Memon, M. H., & Park, S. H. (2019). Feature Selection Based on L1-Norm Support Vector Machine and Effective Recognition System for Parkinson's Disease Using Voice Recordings. IEEE Access, 7, 37718-37734. doi: 10.1109/ACCESS.2019.2906350.
 - [5] Souza, R., Varol, S., & Rueckert, D. (2024). Identifying Biases in a Multicenter MRI Database for Parkinson's Disease Classification: Is the Disease Classifier a Secret Site Classifier? IEEE Journal of Biomedical and Health Informatics, 28(4), 2047-2054. doi: 10.1109/JBHI.2024.3352513.
 - [6] Naz, S., Kamran, I., Gul, S., Hadi, F., & Khalifa, F. (2023). Multi-Model Fusion of CNNs for Identification of Parkinson's Disease Using Handwritten Samples. IEEE Access, 11, 135600-135608. doi: 10.1109/ACCESS.2023.3337804.
 - [7] Yousif, M., Al-Khateeb, B., & Garcia-Zapirain, B. (2024). A New Quantum Circuits of Quantum Convolutional Neural Network for X-Ray Images Classification. IEEE Access, 12, 65660-65671. doi: 10.1109/ACCESS.2024.3396411.
 - [8] Cui, X., Wang, Q., & Zhang, M. (2024). A Multiscale Hybrid Attention Networks Based on Multiview Images for the Diagnosis of Parkinson's Disease. IEEE Transactions on Instrumentation and Measurement, 73, Art no. 2501011. doi: 10.1109/TIM.2023.3315407.
 - [9] Maheshwari, D., Garcia-Zapirain, B., & Sierra-Sosa, D. (2022). Quantum Machine Learning Applications in the Biomedical Domain: A Systematic Review. IEEE Access, 10, 80463-80484. doi: 10.1109/ACCESS.2022.3195044.
 - [10] Panayides, A. S., et al. (2020). AI in Medical Imaging Informatics: Current Challenges and Future Directions. IEEE Journal of Biomedical and Health Informatics, 24(7), 1837-1857. doi: 10.1109/JBHI.2020.2991043.
 - [11] Qin, Z., Jiang, Z., Chen, J., Hu, C., & Ma, Y. (2019). sEMG-Based Tremor Severity Evaluation for Parkinson's Disease Using a Light-Weight CNN. IEEE Signal Processing Letters, 26(4), 637-641. doi: 10.1109/LSP.2019.2903334.
 - [12] Sekaran, K., Sharma, M., & Arora, A. (2024). A Self-Operational Convolutional Neural Networks With Convergent Cross-Mapping and Its Application in Parkinson's Disease Classification. IEEE Access, 12, 83140-83153. doi: 10.1109/ACCESS.2024.3412808.
 - [13] Pesah, Arthur & Cerezo, M. & Wang, Samson & Volkoff, Tyler & Sornborger, Andrew & Coles, Patrick. (2020). Absence of Barren Plateaus in Quantum Convolutional Neural Networks. 10.48550/arXiv.2011.02966.
 - [14] S. Vaghasiya, "NTUA Parkinson Dataset," Kaggle, 2024. [Online]. Available: <https://www.kaggle.com/datasets/shayalvaghasiya/ntua-prakinson>.
 - [15] Chen, S., Ma, K., & Zheng, Y. Med3D: Transfer Learning for 3D Medical Image Analysis. *arXiv preprint arXiv:1904.00625*, 2019.
 - [16] M. J. Cardoso, W. Li, R. Brown, N. Ma, E. Kerfoot, Y. Wang, B. Murray, A. Myronenko, C. Zhao, D. Yang, V. Nath, Y. He, Z. Xu, A. Hatamizadeh, W. Zhu, Y. Liu, M. Zheng, Y. Tang, I. Yang, M. Zephyr, B. Hashemian, S. Alle, M. Z. Darestani, C. Budd, M. Modat, T. Vercauteren, G. Wang, Y. Li, Y. Hu, Y. Fu, B. Gorman, H. Johnson, B. Genereaux, B. S. Erdal, V. Gupta, A. Diaz-Pinto, A. Dourson, L. Maier-Hein, P. F. Jaeger, M. Baumgartner, J. Kalpathy-Cramer, M. Flores, J. Kirby, L. A. D. Cooper, H. R. Roth, D. Xu, D. Bericat, R. Floca, S. K. Zhou, H. Shuaib, K. Farahani, K. H. Maier-Hein, S. Aylward, P. Dogra, S. Ourselin, and A. Feng. MONAI: An open-source framework for deep learning in healthcare. *arXiv preprint arXiv:2211.02701*, 2022. Available: <https://doi.org/10.48550/arXiv.2211.02701>.
 - [17] Hinton, G. E., Srivastava, N., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. R. (2012). *Improving neural networks by preventing co-adaptation of feature detectors*. arXiv preprint arXiv:1207.0580.
 - [18] M. Eapen and J. C. Gore, "Identifying the Functional Architecture of the Human Ventral Tegmental Area and the Substantia Nigra using High Resolution Magnetic Resonance Imaging," in *Proceedings of [relevant conference, if known]*, 2009.
 - [19] Powers, D. M. (2011). *Evaluation: From precision, recall and F-measure to ROC, informedness, markedness and correlation*. Journal of Machine Learning Technologies, 2(1), 37-63.
 - [20] Goodfellow, I., Bengio, Y., & Courville, A. (2021). *Deep Learning: Updated Edition*. MIT Press.
 - [21] Reddi, S. J., Kale, S., & Kumar, S. (2020). *On the Convergence of Adam and Beyond*. Journal of Machine Learning Research, 21(1), 1-31.
 - [22] Tan, M., & Le, Q. V. (2021). *EfficientNetV2: Smaller Models and Faster Training*. In Proceedings of the International Conference on Machine Learning (ICML).
 - [23] Ramachandran, P., Zoph, B., & Le, Q. V. (2020). *Searching for Activation Functions*. IEEE Transactions on Neural Networks and Learning Systems, 31(11), 4589-4603.
 - [24] Cerezo, M., Arrasmith, A., Babbush, R., Benjamin, S. C., Endo, S., Fujii, K., ... & Coles, P. J. (2021). *Variational Quantum Algorithms*. Nature Reviews Physics, 3(9), 625-644.